

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: df = pd.read_csv('diabetes2.csv')
df.head()
```

```
Out[2]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age
0	6	148	72	35	0	33.6	0.627	51
1	1	85	66	29	0	26.6	0.351	33
2	8	183	64	0	0	23.3	0.672	33
3	1	89	66	23	94	28.1	0.167	41
4	0	137	40	35	168	43.1	2.288	33

```
In [3]: df.shape
```

```
Out[3]: (768, 9)
```


```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies           768 non-null    int64
1   Glucose               768 non-null    int64
2   BloodPressure         768 non-null    int64
3   SkinThickness         768 non-null    int64
4   Insulin               768 non-null    int64
5   BMI                   768 non-null    float64
6   DiabetesPedigreeFunction 768 non-null    float64
7   Age                   768 non-null    int64
8   Outcome               768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

There are no null values in the data, so we dont need to remove any rows

```
In [5]: sns.heatmap(df.corr())
```

```
Out[5]: <AxesSubplot:>
```



We can see that the data is not highly correlated, so there is no need to drop any columns

Standardising the input data using StandardScaler

```
In [6]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
df[df.columns[:-1]] = scaler.fit_transform(df[df.columns[:-1]])
df.head()
```

Out[6]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFu
0	0.639947	0.848324	0.149641	0.907270	-0.692891	0.204013	0.4
1	-0.844885	-1.123396	-0.160546	0.530902	-0.692891	-0.684422	-0.5
2	1.233880	1.943724	-0.263941	-1.288212	-0.692891	-1.103255	0.6
3	-0.844885	-0.998208	-0.160546	0.154533	0.123302	-0.494043	-0.9
4	-1.141852	0.504055	-1.504687	0.907270	0.765836	1.409746	5.4

Splitting data into train validation and test sets with seed 0

```
In [7]: train = df.sample(frac=0.7, random_state=0)
val = df.drop(train.index).sample(frac=0.666, random_state=0)
test = df.drop(np.concatenate((train.index, val.index)))

print('Training set size:', train.shape[0])
print('Validation set size:', val.shape[0])
print('Testing set size:', test.shape[0])
```

Training set size: 538
 Validation set size: 153
 Testing set size: 77

Splitting data into input and output

```
In [8]: X_train = train.iloc[:, :-1]
y_train = train.iloc[:, -1]
X_val = val.iloc[:, :-1]
y_val = val.iloc[:, -1]
X_test = test.iloc[:, :-1]
y_test = test.iloc[:, -1]
```

Part 1

Performing Logistic regression

```

In [9]: def sigmoid(x):
        return 1.0 / (1 + np.exp(-x))

class LogisticRegression():
    def __init__(self, alpha, iters):
        self.alpha = alpha
        self.iters = iters

    def bgd(self, x, y, resume):
        # number of training samples and number of features
        m, n = x.shape
        # initialising weights and bias
        if(not resume):
            self.w = np.zeros(n)
            self.b = 0

        self.losses = []
        # iterating until error is negligible or max number of iterations is reached
        for i in range(self.iters):
            self.update(x, y, 'bgd')
            hx = sigmoid(x.dot(self.w) + self.b)
            self.losses.append(loss(hx, y))
            if(resume):
                if(i>2):
                    if(self.losses[-1]>self.losses[-2]):
                        print('early stopping')
                        break

    def sgd(self, x, y, resume):
        # number of training samples and number of features
        m, n = x.shape
        # initialising weights and bias
        if(not resume):
            self.w = np.zeros(n)
            self.b = 0

        self.losses = []
        # iterating until error is negligible or max number of iterations is reached
        for i in range(self.iters):
            for j in range(0, m):
                self.update(x.iloc[j], y.iloc[j], 'sgd')
            hx = sigmoid(x.dot(self.w) + self.b)
            self.losses.append(loss(hx, y))
            if(resume):
                if(i>2):
                    if(self.losses[-1]>self.losses[-2]):
                        print('early stopping')
                        break

    def update(self, x, y, opt):
        x = np.array(x)
        y = np.array(y)
        m = len(x)
        hx = sigmoid(x.dot(self.w) + self.b)
        # gradients for weights and bias
        dw = ( (x.T).dot(hx - y) ) * (1/m)

```

```

db = ( np.sum(hx - y) ) * (1/m)

# updating the weights
self.w = self.w - self.alpha * dw
self.b = self.b - self.alpha * db

def predict(self, x):
    hx = sigmoid(x.dot(self.w) + self.b)
    y_pred = np.where(hx > 0.5, 1, 0)
    return y_pred
# class ends here

def loss(y_pred, y):
    # adding small value so log does not become 0
    epsilon = 1e-10
    # calculates loss for the given predictions
    loss = -np.mean(y*np.log(y_pred + epsilon) + (1-y)*np.log(1-y_pred + epsilon))
    return loss

```

A

BGD loss vs iterations

On training set

```


In [10]: # lr = 0.1, iters = 500
model = LogisticRegression(0.1, 500)
model.bgd(X_train, y_train, False)

```

```

In [11]: plt.plot(np.arange(1, model.iters+1), model.losses)
plt.xlabel('number of epochs')
plt.ylabel('log loss')
plt.show()

```




```
In [12]: accuracy = np.sum(model.predict(X_train)==y_train) / len(y_train)
print('Our model correctly predicts outcome', accuracy*100, '% of the time on the
```

Our model correctly predicts outcome 79.182156133829 % of the time on the training set

On validation set

```
In [13]: # lr = 0.1, iters = 500
model.bgd(X_val, y_val, True)
```

```
In [14]: plt.plot(np.arange(1, model.iters+1), model.losses)
plt.xlabel('number of epochs')
plt.ylabel('log loss')
plt.show()
```



```
In [15]: accuracy = np.sum(model.predict(X_train)==y_train) / len(y_train)
print('Our model correctly predicts outcome', accuracy*100, '% of the time on the
```

Our model correctly predicts outcome 77.88104089219331 % of the time on the validation set

On both the sets it takes around 200 epochs for our model to converge


The accuracy is lower and final loss is higher for validation set as the number of samples is lower in validation set

SGD loss vs iterations

On training set

```
In [16]: # lr = 0.1, iters = 25
model = LogisticRegression(0.1, 25)
model.sgd(X_train, y_train, False)
```

```
In [17]: plt.plot(np.arange(1, model.iters+1), model.losses)
plt.xlabel('number of epochs')
plt.ylabel('log loss')
plt.show()
```




```
In [18]: accuracy = np.sum(model.predict(X_train)==y_train) / len(y_train)
print('Our model correctly predicts outcome', accuracy*100, '% of the time on the
```

Our model correctly predicts outcome 78.62453531598513 % of the time on the training set

On validation set

```
In [19]: # lr = 0.1, iters = 25
model.sgd(X_val, y_val, True)
```

```
In [20]: plt.plot(np.arange(1, model.its+1), model.losses)
plt.xlabel('number of epochs')
plt.ylabel('log loss')
plt.show()
```



```
In [21]: accuracy = np.sum(model.predict(X_train)==y_train) / len(y_train)
print('Our model correctly predicts outcome', accuracy*100, '% of the time on the
```

Our model correctly predicts outcome 77.69516728624535 % of the time on the validation set

On both the sets it takes around 10 epochs for our model to converge

SGD converges in much lesser number of epochs than BGD since the weights are updated with each input point

B


BGD


```
In [22]: def graph(model):
    plt.plot(np.arange(1, model.epochs+1), model.losses)
    plt.xlabel('number of epochs')
    plt.ylabel('log loss')
    plt.show()

def accuracy(model):
    accuracy = np.sum(model.predict(X_test)==y_test) / len(y_test)
    print('Our model correctly predicts outcome', accuracy*100, '% of the time on the test set')


# iters = 2000
lrs = [0.01, 0.0001, 10]
for lr in lrs:
    model = LogisticRegression(lr, 2000)
    model.bgd(X_train, y_train, False)
    print('\n\nLearning rate =', lr)
    graph(model)
    accuracy(model)
```

Learning rate = 0.01




Our model correctly predicts outcome 72.727272727273 % of the time on the test set

Learning rate = 0.0001



Our model correctly predicts outcome 75.32467532467533 % of the time on the testing set

Learning rate = 10



Our model correctly predicts outcome 67.53246753246754 % of the time on the testing set

0.01 learning rate is slower to converge than the original 0.1 taken by us, but it does reach convergence after a few thousand iterations

0.0001 is too low for the learning rate, it will take an unfeasibly large number of iterations to converge.

10 is too high, after some iterations the loss starts bouncing between 2 high loss values and does not reduce any further, it cannot reduce the loss to an acceptable level


SGD

```
In [23]: def graph(model):
    plt.plot(np.arange(1, model.its+1), model.losses)
    plt.xlabel('number of epochs')
    plt.ylabel('log loss')
    plt.show()

def accuracy(model):
    accuracy = np.sum(model.predict(X_test)==y_test) / len(y_test)
    print('Our model correctly predicts outcome', accuracy*100, '% of the time on the test set')


# its = 50
lrs = [0.01, 0.0001, 10]
for lr in lrs:
    model = LogisticRegression(lr, 50)
    model.sgd(X_train, y_train, False)
    print('\n\nLearning rate =', lr)
    graph(model)
    accuracy(model)
```

Learning rate = 0.01




Our model correctly predicts outcome 72.72727272727273 % of the time on the test set

Learning rate = 0.0001



Our model correctly predicts outcome 75.32467532467533 % of the time on the testing set

Learning rate = 10



Our model correctly predicts outcome 71.42857142857143 % of the time on the testing set

0.01 learning rate is slower to converge than the original 0.1 taken by us, but it does reach convergence after about 30 iterations

0.0001 is too low for the learning rate, it will take an infeasibly large number of iterations to converge.

10 is too high, after some iterations the loss starts bouncing between 2 high loss values and does not reduce any further, it cannot reduce the loss to an acceptable level

C

```
In [24]: def confusionMatrix(y, yh):  
         # n is number of classes  
         y = np.array(y)  
         yh = np.array(yh)  
         n = len(np.unique(y))  
         mat = np.zeros((n, n))  
         for i in range(len(y)):  
             mat[y[i]][yh[i]] += 1  
         return mat
```

BGD

```
In [25]: # lr = 0.05, iters = 500  
model = LogisticRegression(0.05, 500)  
model.bgd(X_train, y_train, False)
```

```
In [26]: cmat = confusionMatrix(np.array(y_test), model.predict(X_test))  
print('Confusion matrix:\n', cmat)
```

Confusion matrix:
[[42. 11.]
 [10. 14.]]

```
In [27]: accuracy = (cmat[0][0] + cmat[1][1]) / np.sum(cmat)  
precision = (cmat[1][1]) / (cmat[1][1] + cmat[0][1])  
recall = (cmat[1][1]) / (cmat[1][1] + cmat[1][0])  
f1score = 2* (precision * recall) / (precision + recall)  
  
print('Accuracy:', accuracy)  
print('Precision:', precision)  
print('Recall:', recall)  
print('F1score:', f1score)
```

Accuracy: 0.7272727272727273
Precision: 0.56
Recall: 0.5833333333333334
F1score: 0.5714285714285714

SGD

```
In [28]: # lr = 0.05, iters = 25  
model = LogisticRegression(0.05, 25)  
model.sgd(X_train, y_train, False)
```

```
In [29]: cmat = confusionMatrix(np.array(y_test), model.predict(X_test))
print('Confusion matrix:\n', cmat)
```

```
Confusion matrix:
[[42. 11.]
 [10. 14.]]
```

```
In [30]: accuracy = (cmat[0][0] + cmat[1][1]) / np.sum(cmat)
precision = (cmat[1][1]) / (cmat[1][1] + cmat[0][1])
recall = (cmat[1][1]) / (cmat[1][1] + cmat[1][0])
f1score = 2* (precision * recall) / (precision + recall)

print('Accuracy:', accuracy)
print('Precision:', precision)
print('Recall:', recall)
print('F1score:', f1score)
```

```
Accuracy: 0.7272727272727273
Precision: 0.56
Recall: 0.5833333333333334
F1score: 0.5714285714285714
```


Part 2

0.01 is chosen as the learning rate, with maximum 50 iterations

```
In [31]: from sklearn.linear_model import SGDClassifier
from io import StringIO
import sys
```

A

```
In [32]: old_stdout = sys.stdout
sys.stdout = mystdout = StringIO()
model = SGDClassifier(learning_rate='constant', eta0=0.01, loss='log', max_iter=50)
model.fit(X_train, y_train)
sys.stdout = old_stdout
loss_history = mystdout.getvalue()
loss_list = []
for line in loss_history.split('\n'):
    if(len(line.split("loss: ")) == 1):
        continue
    loss_list.append(float(line.split("loss: ")[-1]))
plt.figure()
plt.plot(np.arange(len(loss_list)), loss_list)
plt.xlabel("Number of epochs")
plt.ylabel("Log Loss")
plt.show()
plt.close()
```



SGDClassifier and our model both have the same kind of plot, and the final loss value is also similar for both

B

```
In [33]: print('SGDClassifier converged in', model.n_iter_, 'iterations')
```

SGDClassifier converged in 11 iterations

sklearn SGDClassifier converges in around 10 epochs while our algorithm took more than 25 epochs, this may be because the sklearn model may use some extra optimisations that we have not

C


```
In [34]: cmat = confusionMatrix(y_test, model.predict(X_test))
accuracy = (cmat[0][0] + cmat[1][1]) / np.sum(cmat)
precision = (cmat[1][1]) / (cmat[1][1] + cmat[0][1])
recall = (cmat[1][1]) / (cmat[1][1] + cmat[1][0])
f1score = 2* (precision * recall) / (precision + recall)

print('Accuracy:', accuracy)
print('Precision:', precision)
print('Recall:', recall)
print('F1score:', f1score)
```

Accuracy: 0.7272727272727273

Precision: 0.56

Recall: 0.5833333333333334

F1score: 0.5714285714285714

Confusion matrix for both the SGDClassifier and our model is the same, so both perform similarly in terms of accuracy, precision, recall and f1score

In []: